

DLCV 109-1

HW2

陳苑彣 r08942085

Problem 1

1-1

Backbone: pertained vgg16

Layer: 將兩次 upsample 的結果 concat，最後過三層 fully connected layers。

```
Vggfcn(  
    (pretrain): VGG(  
        (features): Sequential(  
            (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
            (1): ReLU(inplace=True)  
            (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
            (3): ReLU(inplace=True)  
            (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
            (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
            (6): ReLU(inplace=True)  
            (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
            (8): ReLU(inplace=True)  
            (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
            (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
            (11): ReLU(inplace=True)  
            (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
            (13): ReLU(inplace=True)  
            (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
            (15): ReLU(inplace=True)  
            (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
            (17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
            (18): ReLU(inplace=True)  
            (19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
            (20): ReLU(inplace=True)  
            (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
            (22): ReLU(inplace=True)  
            (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
            (24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
            (25): ReLU(inplace=True)  
            (26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
            (27): ReLU(inplace=True)  
            (28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
            (29): ReLU(inplace=True)  
            (30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
        )  
        (avgpool): AdaptiveAvgPool2d(output_size=(7, 7))  
        (classifier): Sequential(  
            (0): Linear(in_features=25088, out_features=4096, bias=True)  
            (1): ReLU(inplace=True)  
            (2): Dropout(p=0.5, inplace=False)  
            (3): Linear(in_features=4096, out_features=4096, bias=True)  
            (4): ReLU(inplace=True)  
            (5): Dropout(p=0.5, inplace=False)  
            (6): Linear(in_features=4096, out_features=1000, bias=True)  
        )  
    )  
  
(layer1): Sequential(  
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): ReLU(inplace=True)  
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (3): ReLU(inplace=True)  
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (6): ReLU(inplace=True)  
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (8): ReLU(inplace=True)  
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (11): ReLU(inplace=True)  
    (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (13): ReLU(inplace=True)  
    (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (15): ReLU(inplace=True)  
    (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (18): ReLU(inplace=True)  
    (19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (20): ReLU(inplace=True)  
    (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (22): ReLU(inplace=True)  
    (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (25): ReLU(inplace=True)  
    (26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (27): ReLU(inplace=True)  
    (28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (29): ReLU(inplace=True)  
    (30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
)  
  
(conv1): Conv2d(512, 512, kernel_size=(1, 1), stride=(1, 1))  
(up): Upsample(scale_factor=2.0, mode=bilinear)  
(conv2): Conv2d(512, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
(conv3): Conv2d(256, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
(fc): Sequential(  
    (0): Linear(in_features=6144, out_features=2048, bias=True)  
    (1): Linear(in_features=2048, out_features=256, bias=True)  
    (2): Linear(in_features=256, out_features=50, bias=True)  
)  
)
```

```
38 class Vggfcn(nn.Module):  
39     def __init__(self, pretrain=False):  
40         super(Vggfcn, self).__init__()  
41         self.pretrain = pretrain  
42         self.layer1 = pretrain.features  
43         self.conv1 = nn.Conv2d(512, 512, 1)  
44         self.up = nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True)  
45         self.conv2 = nn.Conv2d(512, 256, 3, stride=1, padding=1)  
46         self.conv3 = nn.Conv2d(256, 128, 3, stride=1, padding=1)  
47         self.fc = nn.Sequential(  
48             nn.Linear(6144, 2048),  
49             nn.Linear(2048, 256),  
50             nn.Linear(256, 50),  
51         )  
52  
53     def forward(self, input):  
54         x = input  
55         x = self.layer1(x) #[32, 512, 1, 1]  
56         x = self.conv1(x)  
57         x = self.up(x) # [512, 2, 2]  
58         x = self.conv2(x)  
59         x = self.up(x) # [256, 4, 4]  
60         x2 = self.conv3(x)  
61         x3 = torch.cat([x, x2], dim=1)  
62         x3 = x3.view(x3.shape[0], -1)  
63         output = self.fc(x3)  
64  
65         return output
```

Backbone

Backbone

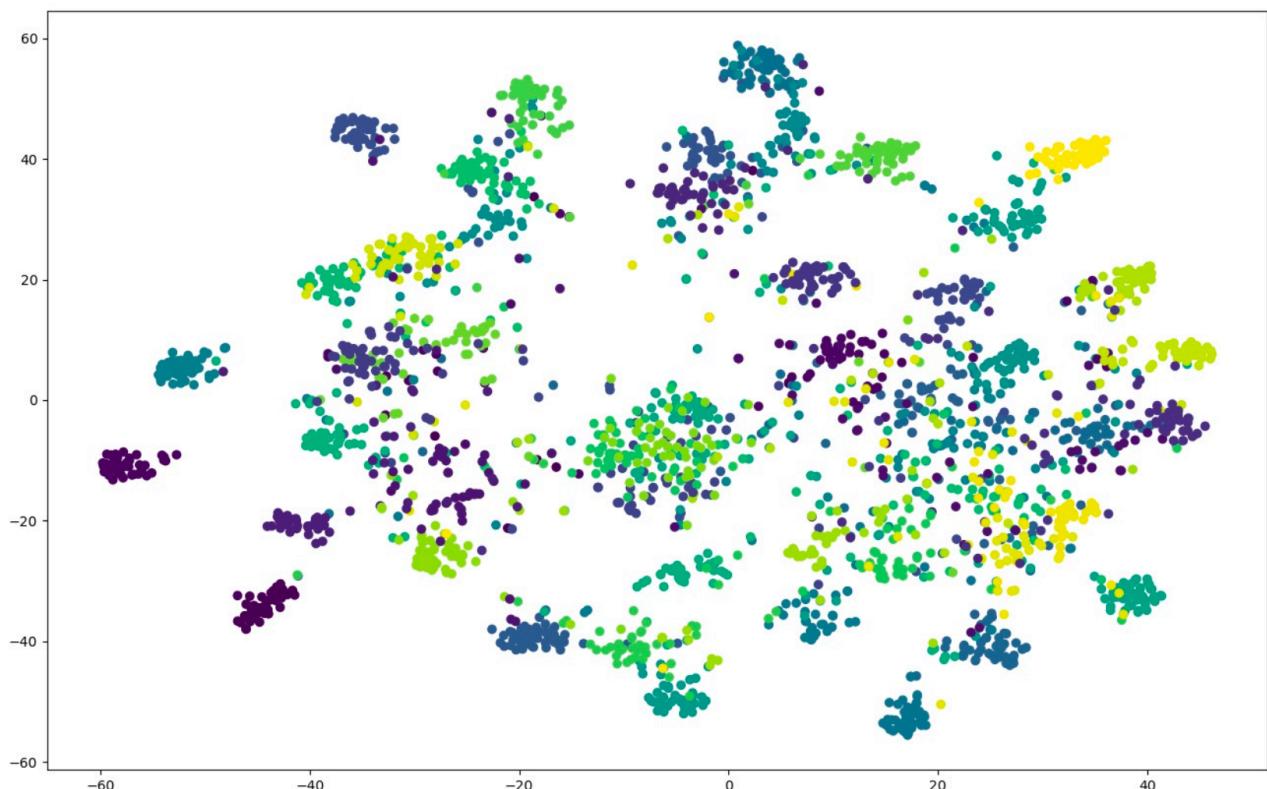
1-2

Accuracy on validation set:

Accuracy: 1821/2500 (73%)

1-3

Tsne result: 可以看到在 accuracy 73% 的分類結果下，資料點在高維分佈降維後的分群結果，有些有明顯的分群（例如周圍的部份），大致有 30 群，但如中間或中右的資料點，模型就沒有分群分得很好。



Problem2

2-1

Backbone: vgg16 (下圖中的 layer1)

Layer: 使用 convTranspose 五次，每次圖片放大兩倍，並且有經過 batch normalize 及 relu。

reference: <https://github.com/pochih/FCN-pytorch/blob/master/python/fcn.py>

```
Vggfcn32C
(features): VGG(
    (pretrain): Sequential(
        (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): ReLU(inplace=True)
        (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (3): ReLU(inplace=True)
        (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (6): ReLU(inplace=True)
        (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (8): ReLU(inplace=True)
        (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (11): ReLU(inplace=True)
        (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (13): ReLU(inplace=True)
        (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (15): ReLU(inplace=True)
        (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (18): ReLU(inplace=True)
        (19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (20): ReLU(inplace=True)
        (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (22): ReLU(inplace=True)
        (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (25): ReLU(inplace=True)
        (26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (27): ReLU(inplace=True)
        (28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (29): ReLU(inplace=True)
        (30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (avgpool): AdaptiveAvgPool2d(output_size=(7, 7))
    (classifier): Sequential(
        (0): Linear(in_features=25088, out_features=4096, bias=True)
        (1): ReLU(inplace=True)
        (2): Dropout(p=0.5, inplace=False)
        (3): Linear(in_features=4096, out_features=4096, bias=True)
        (4): ReLU(inplace=True)
        (5): Dropout(p=0.5, inplace=False)
        (6): Linear(in_features=4096, out_features=1000, bias=True)
    )
)
layer1: Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace=True)
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU(inplace=True)
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (6): ReLU(inplace=True)
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (8): ReLU(inplace=True)
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace=True)
    (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (13): ReLU(inplace=True)
    (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (15): ReLU(inplace=True)
    (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (18): ReLU(inplace=True)
    (19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (20): ReLU(inplace=True)
    (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (22): ReLU(inplace=True)
    (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (25): ReLU(inplace=True)
    (26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (27): ReLU(inplace=True)
    (28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (29): ReLU(inplace=True)
    (30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(relu): ReLU(inplace=True)
(dev1): ConvTranspose2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), output_padding=(1, 1))
(dev2): ConvTranspose2d(512, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), output_padding=(1, 1))
(dev3): ConvTranspose2d(256, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), output_padding=(1, 1))
(dev4): ConvTranspose2d(128, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), output_padding=(1, 1))
(dev5): ConvTranspose2d(64, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), output_padding=(1, 1))
(bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(bn4): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(bn5): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(conv7): Conv2d(32, 7, kernel_size=(1, 1), stride=(1, 1))
```

```
class Vggfcn32(nn.Module):
    def __init__(self, pretrain=False):
        super(Vggfcn32, self).__init__()
        self.pretrain = pretrain
        self.layer1 = pretrain.features
        self.relu = nn.ReLU(inplace=True)
        self.dev1 = nn.ConvTranspose2d(512, 512, kernel_size=3, stride=2, padding=1, dilation=1, output_padding=1)
        self.bn2 = nn.BatchNorm2d(256)
        self.bn3 = nn.BatchNorm2d(128)
        self.bn4 = nn.BatchNorm2d(64)
        self.bn5 = nn.BatchNorm2d(32)
        self.conv7 = nn.Conv2d(32, 7, 1)

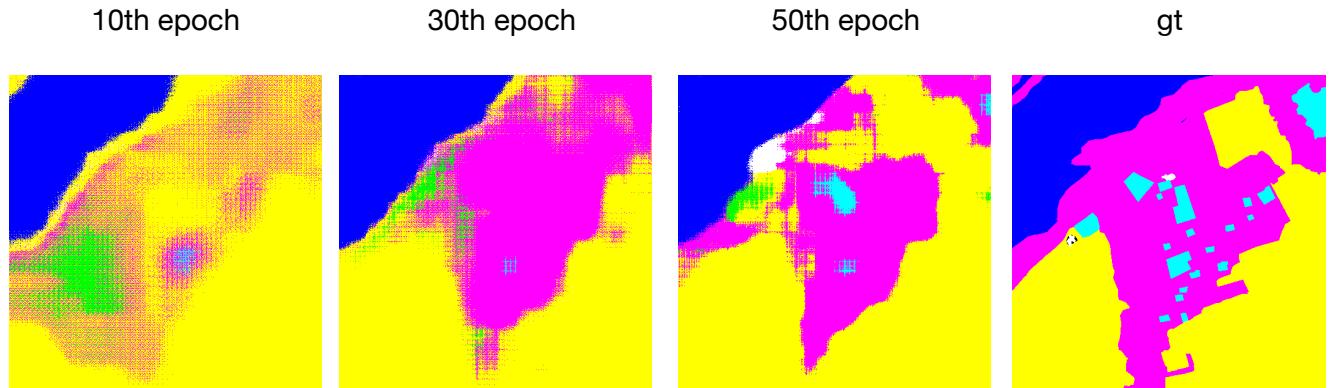
    def forward(self, input):
        x = input
        x = self.layer1(x)
        x = self.bn1(self.relu(self.dev1(x)))
        x = self.bn2(self.relu(self.dev2(x)))
        x = self.bn3(self.relu(self.dev3(x)))
        x = self.bn4(self.relu(self.dev4(x)))
        x = self.bn5(self.relu(self.dev5(x)))
        x = self.conv7(x)
        # pdb.set_trace()
        # x = self.convTranspose(x)
        # x = F.upsample_bilinear(x, input.size()[2:])
        return x
```

Backbone

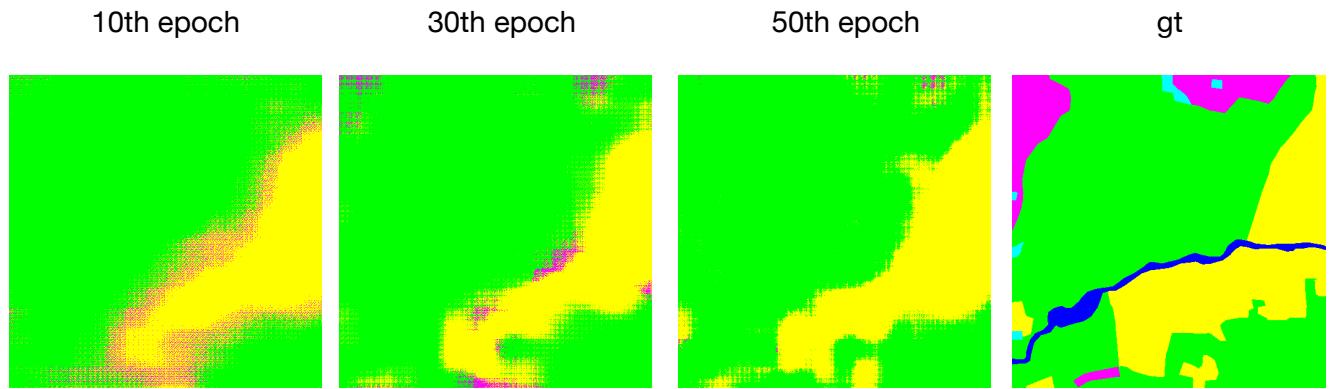
Backbone

2-2

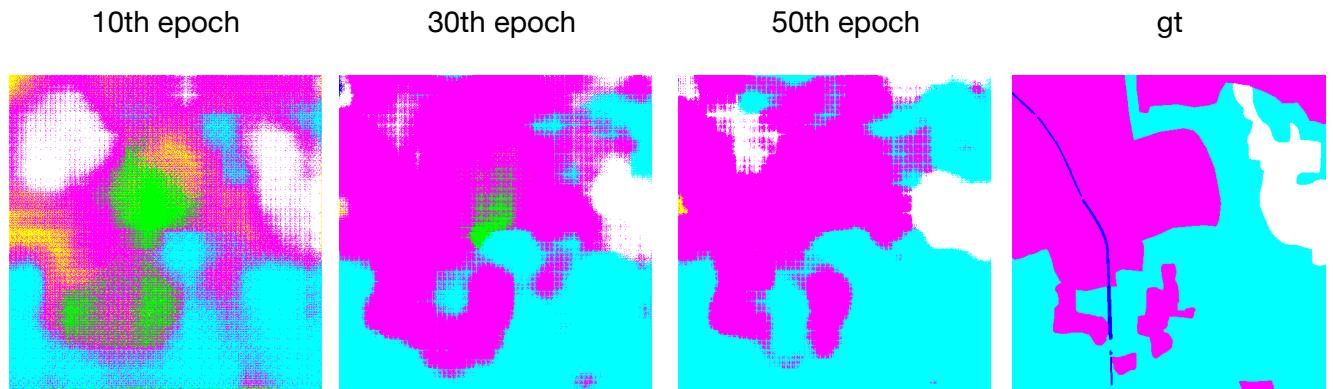
[0010_sat.jpg]



[0097_sat.jpg]



[107_sat.jpg]



2-3

Backbone: vgg16_bn

Layer: 模仿 fcn8s 的做法，把 vgg16_bn 前面 encode 的 features 與後來 upsample 的 features pixel wise 加起來，必經過幾次 convolution。本次 convTranspose 有先 initialize 成 bilinear upsample。

Reference: <https://github.com/wkentaro/pytorch-fcn/blob/master/torchfcn/models/fcn32s.py>

```

(block1): Sequential(
(0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(2): ReLU(inplace=True)
(3): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(4): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(5): ReLU(inplace=True)
(6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(block2): Sequential(
(7): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(8): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(9): ReLU(inplace=True)
(10): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(11): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(12): ReLU(inplace=True)
(13): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(block3): Sequential(
(14): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(15): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(16): ReLU(inplace=True)
(17): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(18): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(19): ReLU(inplace=True)
(20): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(21): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(22): ReLU(inplace=True)
(23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(block4): Sequential(
(24): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(25): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(26): ReLU(inplace=True)
(27): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(28): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(29): ReLU(inplace=True)
(30): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(31): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(32): ReLU(inplace=True)
(33): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(block5): Sequential(
(34): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(35): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(36): ReLU(inplace=True)
(37): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(38): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(39): ReLU(inplace=True)
(40): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(41): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(42): ReLU(inplace=True)
(43): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
)
(relu): ReLU(inplace=True)
(dev1): ConvTranspose2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), output_padding=(1, 1))
(dev1_1): ConvTranspose2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), output_padding=(1, 1))
(dev2): ConvTranspose2d(256, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), output_padding=(1, 1))
(dev3): ConvTranspose2d(128, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), output_padding=(1, 1))
(dev4): ConvTranspose2d(64, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), output_padding=(1, 1))
(bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(bn1_1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(bn4): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(bn5): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(conv7): Conv2d(32, 7, kernel_size=(1, 1), stride=(1, 1))
(conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(conv2): Conv2d(512, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(conv3): Conv2d(256, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(conv4): Conv2d(128, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(conv5): Conv2d(64, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(conv6): Conv2d(32, 7, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
)

```

Backbone

```

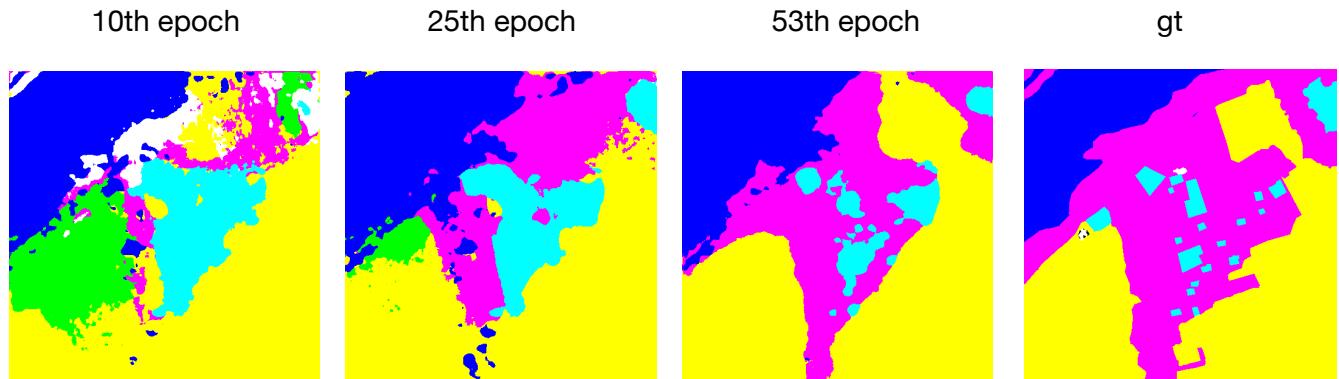
146 class Vggfcn8(nn.Module):
147     def __init__(self, pretrain=False):
148         super(Vggfcn8, self).__init__()
149         print("It's Vggfcn8_bn")
150         self.pretrain = pretrain
151         self.pretrain_all = pretrain.features
152
153         self.block1 = pretrain.features[0:7] #[0:5]
154         self.block2 = pretrain.features[7:14] #[5:10]
155         self.block3 = pretrain.features[14:24] #[10:17]
156         self.block4 = pretrain.features[24:34] #[17:24]
157         self.block5 = pretrain.features[34:44] #[24:31]
158
159         self.relu = nn.ReLU(inplace=True)
160         self.dev1 = nn.ConvTranspose2d(512, 512, kernel_size=3,
161                                     stride=2, padding=1, dilation=1, output_padding=1)
162         self.dev1_1 = nn.ConvTranspose2d(512, 512, kernel_size=3,
163                                     stride=2, padding=1, dilation=1, output_padding=1)
164         self.dev2 = nn.ConvTranspose2d(256, 256, kernel_size=3,
165                                     stride=2, padding=1, dilation=1, output_padding=1)
166         self.dev3 = nn.ConvTranspose2d(128, 128, kernel_size=3,
167                                     stride=2, padding=1, dilation=1, output_padding=1)
168         self.dev4 = nn.ConvTranspose2d(64, 64, kernel_size=3,
169                                     stride=2, padding=1, dilation=1, output_padding=1)
170
171         self.bn1 = nn.BatchNorm2d(512)
172         self.bn1_1 = nn.BatchNorm2d(512)
173         self.bn2 = nn.BatchNorm2d(256)
174         self.bn3 = nn.BatchNorm2d(128)
175         self.bn4 = nn.BatchNorm2d(64)
176         self.bn5 = nn.BatchNorm2d(32)
177         self.conv7 = nn.Conv2d(32, 7, 1)
178
179         self.conv1 = nn.Conv2d(512, 512, 3, stride=1, padding=1)
180         self.conv2 = nn.Conv2d(512, 256, 3, stride=1, padding=1)
181         self.conv3 = nn.Conv2d(256, 128, 3, stride=1, padding=1)
182         self.conv4 = nn.Conv2d(128, 64, 3, stride=1, padding=1)
183         self.conv5 = nn.Conv2d(64, 32, 3, stride=1, padding=1)
184         self.conv6 = nn.Conv2d(32, 7, 3, stride=1, padding=1)
185
186         self._initialize_weights()
187     def _initialize_weights(self):
188         for m in self.modules():
189             if isinstance(m, nn.ConvTranspose2d):
190                 assert m.kernel_size[0] == m.kernel_size[1]
191                 initial_weight = get_upsampling_weight(
192                     m.in_channels, m.out_channels, m.kernel_size[0])
193                 m.weight.data.copy_(initial_weight)
194
195     def forward(self, input):
196         x = input
197
198         x1 = self.block1(x) # 64, 256, 256
199         x2 = self.block2(x1) # 128, 128, 128
200         x3 = self.block3(x2) # 256, 64, 64
201         x4 = self.block4(x3) # 512, 32, 32
202         x5 = self.block5(x4) # 512, 16, 16
203
204         x5 = self.conv1(x5)
205         u1 = self.relu(self.bn1(self.dev1(x5))) # 512, 32, 32
206         u2 = self.relu(self.bn1_1(self.dev1_1(u1+x4))) # 256, 64, 64
207         u2 = self.conv2(u2) # 256, 32, 32
208
209         u3 = self.relu(self.bn2(self.dev2(u2+x3))) # 128, 128, 128
210         u3 = self.conv3(u3) # 256, 64, 64
211
212         u4 = self.relu(self.bn3(self.dev3(u3+x2))) # 64, 256, 256
213         u4 = self.conv4(u4) # 128, 128, 128
214
215         u5 = self.relu(self.bn4(self.dev4(u4+x1))) # 32, 512, 512
216         u5 = self.conv5(u5) # 64, 256, 256
217
218         output = self.conv6(u5)
219         # pdb.set_trace()
220
221         return output
222

```

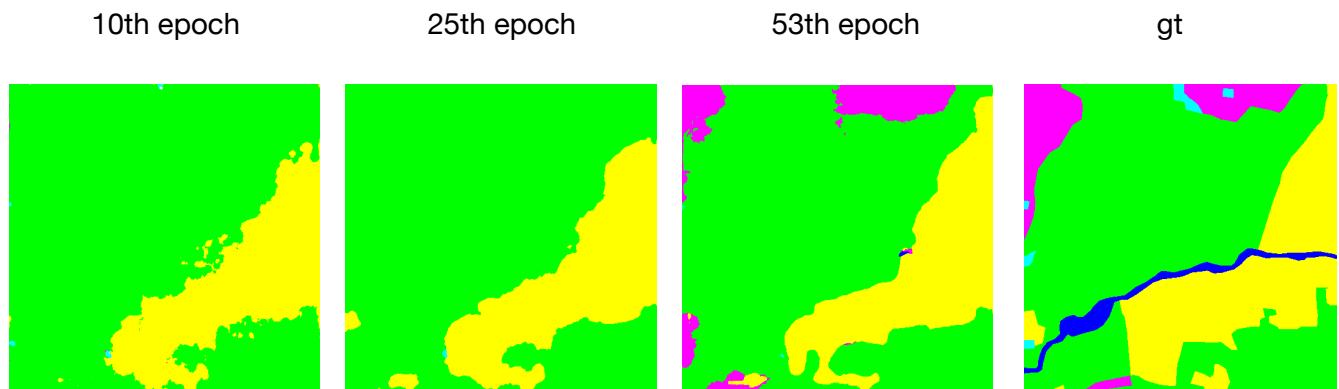
Backbone

2-4

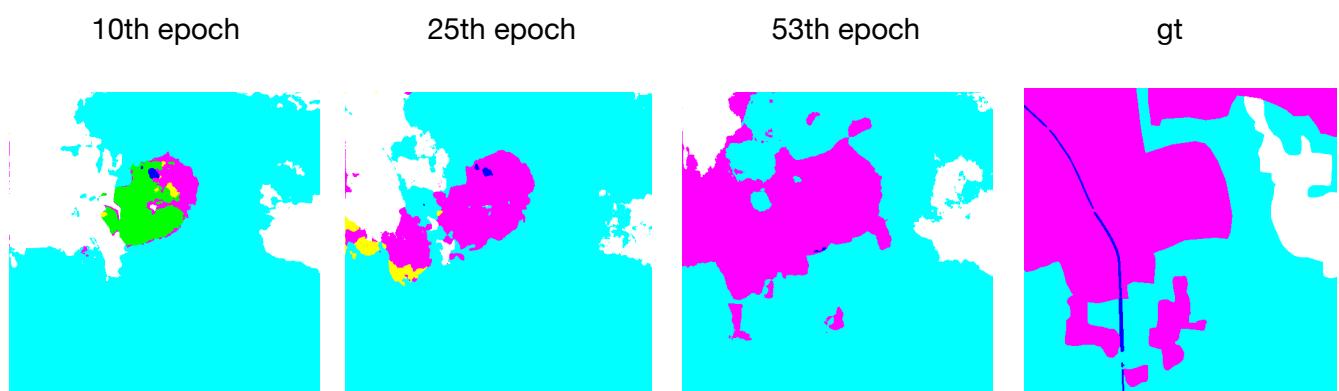
[0010_sat.jpg]



[0097_sat.jpg]



[107_sat.jpg]



2-5

兩個 segmentation model 差異除了 backbone 有沒有加 batch normalize 之外，另外就是後面 upsample 的做法：improvement model 我採用 fcn8s，與 fcn32s 不同的是它結合了 downsample 後的 feature 資訊，因此相較於 fcn32s，fcn8s 可以保留較多原始圖片的 context，因為 downsample 主要是取出 local context 內容資訊，upsample 主要取出 location 位置資訊，因此若

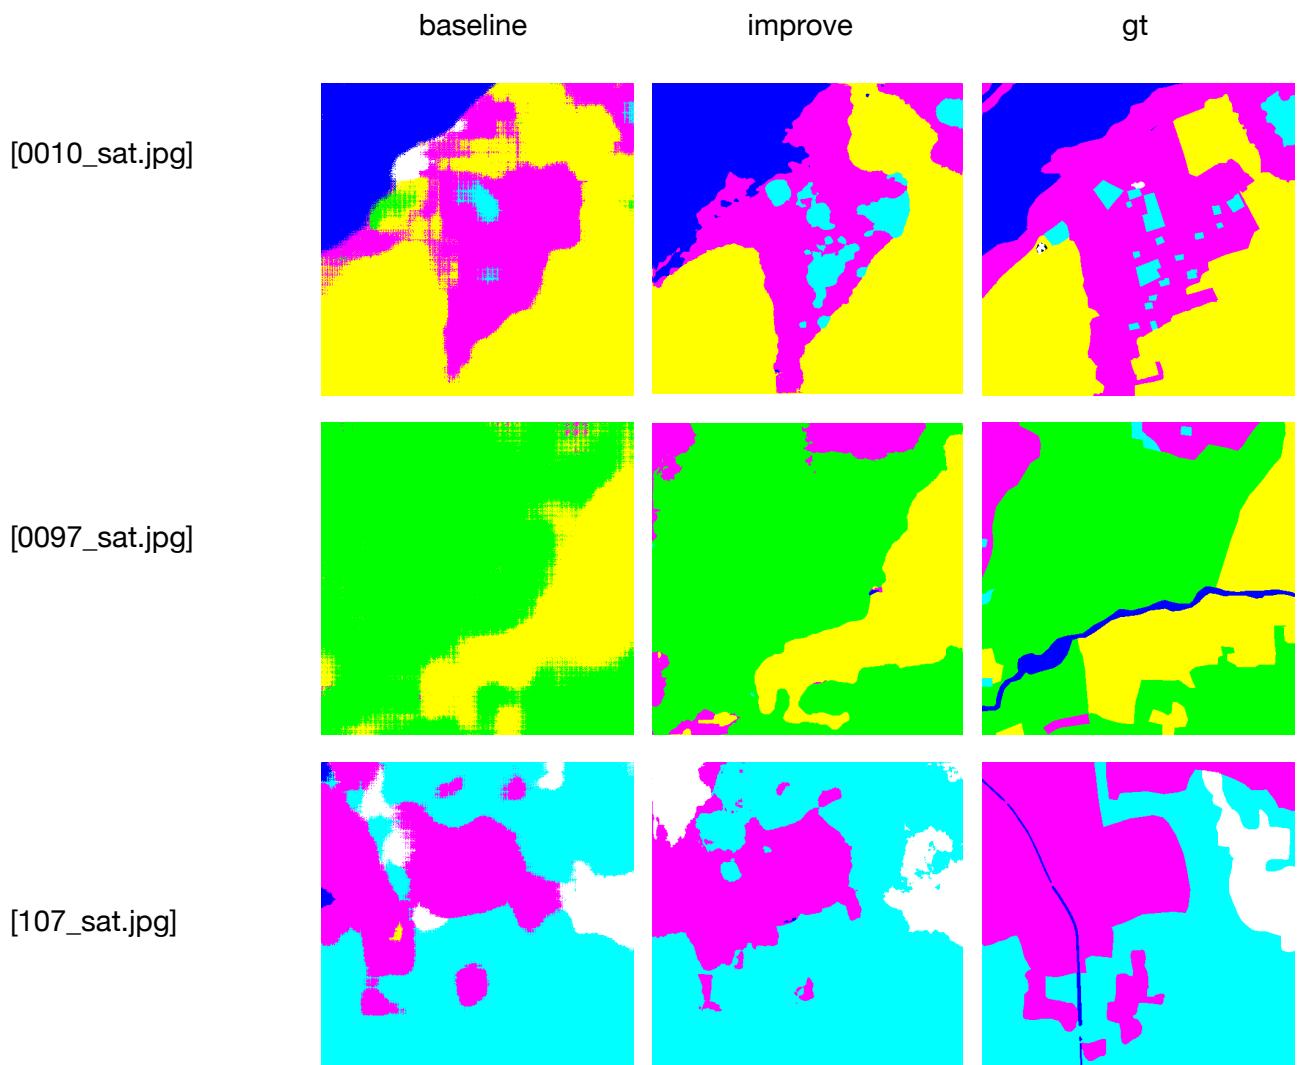
可以用一些 skip connection (concat or pixel wise sum) , 可以有效同時保留 context information 與 spatial information 。

Baseline model 我取第 50 epoch 的 model , validation meanIOU 成績為 67% 。

Improvement model 我取第 53 epoch 的 model , validation meanIOU 成績為 72% 。

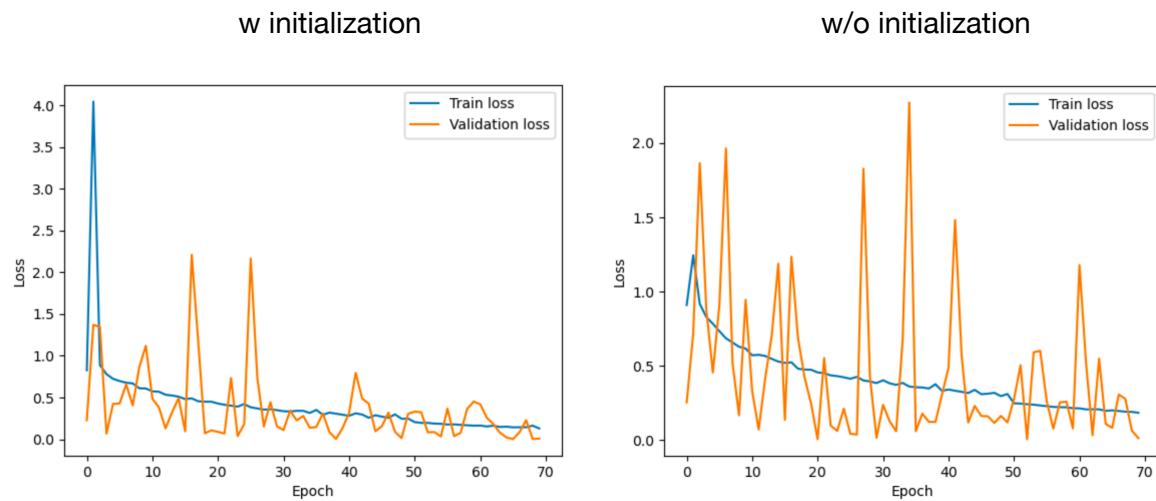
| baseline | improve |
|--------------------|--------------------|
| class #0 : 0.65544 | class #0 : 0.77022 |
| class #1 : 0.87655 | class #1 : 0.89533 |
| class #2 : 0.32324 | class #2 : 0.40155 |
| class #3 : 0.79471 | class #3 : 0.81994 |
| class #4 : 0.73741 | class #4 : 0.77723 |
| class #5 : 0.65466 | class #5 : 0.70049 |
| mean_iou: 0.673666 | mean_iou: 0.727459 |

由下圖可以看到，兩者類別的位置資訊大致上都接近 ground truth ，但 improvement model 的 segmentation 邊緣形狀比較接近 ground truth ，也比較銳利，因此 improvement model 可以有效的保留 context 資訊 。



為了使 meanIOU 成績更好，我有使用 torch.optim 的 lr_scheduled，每經過 50 個 epoch learning rate 就會乘上 0.5 下降，讓整體（train & validate）的 meanIOU 提升，另外 improvement model 中的每一層 convTransport 我都有 initialize 成 bilinear upsample，好處是可以在起始 epoch 就有不錯的成績，應可以更快收斂，另外，為了解決 overfitting，在做 improvement model 時我也有額外做 data augmentation，讓 training data 有 0.5 的機率水平翻轉，以及 0.5 的機率垂直翻轉，讓原本 improvement 的 meanIOU 從 70% 提升為 72%。

下方為有初始化及無初始化的 loss 圖，可以看到沒有初始化的 validation loss 超過 1.5 的次數很多，震盪相對多。



reference: <https://arxiv.org/pdf/1605.06211.pdf>
“no collaborators”